CONSOLIDATED 23 AUG · UPDATED 24 AUG 2026

Resume Builder Backlog

Every planned or specified piece of work that is not yet fully built — swept from `.gemini` working docs, all of `docs/` (plans, specs, audits, idea lists and the review programme), and `IDEAS.md` — deduplicated against the actual code, with the evidence that decided each call.

28	8	0	102	11
OPEN	PARTIAL	STALE DOC	DONE	WON'T DO

152 tracked items across 23 sources · 36 need a decision or a commit — 27 feature, 9 cosmetic, 3 tech debt

Re-verified 23 Aug — the master audit was mostly already closed

An earlier pass of this page marked several master-audit findings open on the strength of truncated greps. Every one was re-checked against the source: **F6, F9, F10, F11, F18 and F22 were already fixed** — including **F11, the only Critical in the whole audit**. `docs/review/master_audit_document.md` now carries a verified status table at the top, so it stops triggering re-triage.

One finding had quietly gone live, and is now fixed. F20 was filed as "aspirational" because mobile Lite Mode didn't exist yet. `requirements-lite.txt` has since shipped and deliberately omits pandas/numpy — while `orchestrator.py` still imported both at module level, and every entry point imports it. A Lite install died on plain `resume` with `ModuleNotFoundError` before printing anything. Both are now lazy across four modules, with `tests/test_lite_mode_imports.py` pinning it.

Two separate F-numbering schemes still exist. The master audit's F1–F22 and the bootstrap remediation spec's own F1–F16 are unrelated lists. Closing one never closed the other.

The three `*unimplemented_plans.md` files are not a backlog. All three carry the identical 1,525 unchecked boxes — three views of one scrape that counted every `- [ ]` in the repo as unfinished. This project never ticks checkboxes, so shipped work reads as pending. They live in Gemini's own session state, so they're yours to delete; this page replaces them.

24 Aug — the Charm cockpit round landed, and a review pass closed 19 findings

Gemini's cockpit expansion (mouse interactivity, a Skills Gap Matrix, action keybindings) is **now committed** in three commits: a `gofmt` sweep of five files that were already unformatted at HEAD, the feature work with its remediations, and the coverage-scoring fix. Both suites green, all gates clean.

Three of the findings were features that shipped invisible. `renderMissionControl()` was fully written and never called; the `scan` / `batch_evaluate` / `sweep_stale` Python subcommands had no Go caller; and the Skills Gap Matrix raised `FileNotFoundError` on database-only jobs, which is most of them. Passing tests said nothing about any of it.

Two findings were tests that reported success while asserting nothing. All four mouse-click tests wrapped their only assertion in `if info != nil`, so a zone that stopped registering read as a pass. Separately, a `time.Sleep` added to `zone.Scan` to mask a race was hiding a real one — removing it made the zone test fail under `-race` immediately.

The Skills Gap Matrix number was measured, not assumed. Against the real 844-bullet corpus, the shipped `(x−0.50)/0.35` rescale pinned 95% of queries at 100% and could never display below 63.9% — a gap matrix structurally unable to show a gap. Coverage is now a percentile rank, needing no hand-tuned constants. Recorded in `CLAUDE.md`.

24 Aug — widened to every plan and spec, not just the ones already being tracked

The register previously covered 14 sources. A sweep of all ~180 markdown files under `docs/` found **five more that hold planned work** and had never been tracked anywhere: `DASHBOARD_IDEAS.md`, the eleven planning and audit documents in `docs/gemini_files/`, and the review programme's own three plan files. They are now in, at **19 sources / 150 items**.

The biggest thing the sweep surfaced is a gate, not an item. `application_log` and `contacts` are both still **0 rows** — confirmed against the live database, not the docs. Roughly half of DASHBOARD_IDEAS is blocked on that one unlock: email status write-back, which `inbox_sync.py` deliberately did not do yet at the time of this sweep. Everything funnel-shaped waits behind it, including the funnel itself, which renders today with nothing in it. *Update 24 Aug: the write-back capability has since shipped behind an opt-in `--apply` flag — see the note below. `application_log` is still 0 rows because it hasn't been run against real mail yet, not because the mechanism is missing.*

Two entries are marked “stale doc” rather than open or done. Four UX audits in `docs/gemini_files/` carry zero resolution markers between them, and the TUI critique carries one. Some of their findings have clearly been fixed since — but nothing records which, so they cannot honestly be called either. They need one triage pass or an explicit retirement; guessing a status would have put fiction in the register.

Research is listed as explicitly not a backlog. ~30 deep-research reports under `morgan_research/` and `FeatureResearch/` read like plans and have twice been mistaken for one. They are inputs to F3/F4/F5 — recorded here so the next sweep does not re-triage them.

24 Aug — four remediation tasks landed and independently verified

All four confirmed against the actual diffs and a live test run, not the handoff notes alone. Doc hygiene (F19/F21) is closed — `refactoring_plan.md:4-11` now carries an Audit Status Banner reconciling it against `CORRECTIONS.md`, and `master_audit_document.md` shows both findings Closed. D11's token-bucket rate limiter is real: thread-safe, 12 RPM default, wired into both `generate()` and `embed()` in `gemini_client.py`. The full suite ran clean at **2,353 tests, 100% passing**, matching the claim exactly.

Two of the four are progress, not closure — marked `PARTIAL`, not done. The cookie-scraping item is *hardened*: `scan_linkedin.py`'s top-level `import browser_cookie3` is gone and extraction is now lazy with real error handling pointing at the safer Playwright login — but `browser_cookie3` is still the mechanism when a user picks it, so the roadmap item (*replace* with safe APIs) isn't actually done. The email write-back gate is the more consequential case: `db.py`, `jd_manager.py` and `inbox_sync.py --apply` genuinely ship a tested write path, correctly guarded by the existing `_is_unisolated_test_write` pattern — but `application_log` was still **0 rows** in the live database at verification time, because it hasn't been run against real mail yet. Half of DASHBOARD_IDEAS stays gated until someone actually runs `--apply`.

One soft-fail pattern worth a conscious look, not a blocker. `jd_manager.save_application_status()` wraps its DB write in a bare `except Exception: logging.warning(...)` — reasonable since the JSON file stays the source of truth and the DB is a secondary index, but it's the kind of degrade-instead-of-fail-closed choice this codebase usually avoids (contrast `gemini_client`'s test-network guard, which raises rather than degrades). *Update 24 Aug: fixed in Sprint 3 — see the note below.*

24 Aug — Sprint 3: tech debt, docs, and a security cleanup, all confirmed

The soft-fail pattern flagged two notes up is fixed. `jd_manager.py`'s blanket `except Exception` blocks in `_sync_jd_to_db()`, `save_application_status()`, `move_jd_to()`, and `archive_job()` now catch specific exception classes (`sqlite3.OperationalError`/`sqlite3.DatabaseError` for DB failures, `json.JSONDecodeError`/`OSError`/`UnicodeDecodeError` for file I/O) — confirmed by direct read, 25+ sites in the file now use narrow catches, with a genuine unexpected-failure fallback logging a full traceback via `exc_info=True` rather than swallowing it. The security item is real too: the operator's `CONTACT_INFO` dict is emptied out at `profiles/morgan/fixed_content.py` (identity now flows purely from `profile.yml` via the existing fill-only mechanism in `profile_paths.py`) — confirmed the fill still produces the correct real values at render time, and both guardrail test files pass (24/24). Note the batch message cites this file as `fixed_content.py` — it actually lives at `profiles/morgan/fixed_content.py`, not `scripts/`, since it was migrated profile-scoped earlier this session; same recurring citation-imprecision pattern, functionally correct regardless. `master_audit_document.md` genuinely reconciles F19/F21 as Closed (F1 correctly still shown as Decomposed/Progress, not overclaimed as done). README/operations/FAQ doc updates spot-checked as real, not just claimed. Full suite: **2,409 tests, 0 failures**, matching the claim. Verified directly (git status, diff read, two live checks, one direct test run) rather than a fork.

24 Aug — Sprint 4: three real debt closures, two thin documentation claims

Three of four tech-debt items are genuinely done. Docker/DevContainer: a real, substantive `Dockerfile` (Python 3.12, Node 20 LTS + Playwright Chromium, Go 1.23.6, Typst, fonts, SQLite3), `.dockerignore`, and `.devcontainer/devcontainer.json` — not stubs, confirmed by reading them. The circular-import item is the one worth calling out: this closes the exact gap flagged after Sprint 2. `menu.py` now has its own `_display_main_banner()` wrapper that explicitly passes `picker.count_active_roles`/`jd_manager.count_completed_resumes`/`picker.count_unevaluated_roles`, and grepping confirms it's the ONLY production call site left — so the lazy `import picker` fallback inside `cli_art._stats_line_text()` is now genuinely unreachable outside tests. Auto-reembedding gained real `needs_reembed()` / `reembed(blocking=)` helpers, confirmed by direct read, 16 new tests. Full suite: **2,417 tests, 0 failures**, matches the claim.

The SQLite-unification item is correctly still marked `PARTIAL`, not `done`. `resume reconcile [--apply] [--profile]` is real (confirmed in `cli.py`) — but it's a new CLI wrapper around the pre-existing `reconcile_jd_status.py` (backup-before-write, dry-run default were already documented in this register before today), not new reconciliation logic. It's discoverability, not the SQLite-as-single-source-of-truth design change the item actually asks for, which stays intentionally undone by design.

The two documentation items don't hold up under a direct check, and are left `open`. "Retired" the three `*unimplemented_plans.md` scrapes was claimed to be recorded in both `implementation_notes_2026_08_24.md` AND `master_audit_document.md` — only the first actually mentions it; grepping `master_audit_document.md` for it finds nothing. Separately, the files themselves were never at risk: they live in `~/.gemini/`, entirely outside this git repo (another tool's global session state, not repo content), so there was never anything here to delete or track — the backlog's own prior note already had this right. The FeatureResearch classification claim has no findable trace anywhere either. Neither is a functional problem, just paperwork that didn't land where claimed — leaving both rows as they were.

24 Aug — the last three, closed with two real caveats worth acting on

Both documentation items are genuinely closed this time. The `master_audit_document.md` citation that was checked and found false last round is real now: DOC-4 and DOC-5 both exist there (confirmed by direct read), and `docs/to_do/FeatureResearch/README.md` genuinely indexes all 5 domains to F3/F4/F5. The `~/.gemini/` scrape files were confirmed still untouched on disk (correctly — they're outside this repo and never should be deleted from here), just formally dropped from tracking. Full suite: **2,418 tests, 0 failures**, matches the claim.

SQLite unification makes real progress but stays `PARTIAL`, and surfaced two things worth fixing, not just noting. `db.py`'s new `get_job_count()` / `get_active_jobs()` / `get_completed_resumes_count()` and `JDTracker.mark_completed()`'s new call to `db.update_job_status()` are all real and test-covered — but neither issue below is exercised by the new test, which only checks the simple case. **First:** that `update_job_status()` call is wrapped in a bare `except Exception: pass` with zero logging — a step backward from the logged-fallback pattern this exact session already established in this same file's other DB writes; a silent sync failure here would be invisible. **Second:** `get_completed_resumes_count()` takes `MAX()` of two independent counts (distinct `application_log` job_ids vs. `jobs` rows in a completed-like status) instead of a union of actual job IDs — if the two sets diverge, which is plausible given they're populated by different code paths (email matching vs. resume tailoring), this silently undercounts. Both are currently dormant since live data has 0 rows in both sources, so nothing is wrong today — but worth fixing before this count is trusted anywhere real.

24 Aug — both flagged bugs fixed and confirmed real

Both issues from the note above are genuinely resolved, confirmed by direct read. `get_completed_resumes_count()` now runs a true SQL `UNION` of job IDs from `application_log` and `jobs` (not the old `MAX()`-of-two-counts) — the exact fix needed, and `test_db.py`'s test was rewritten to actually exercise the divergent-ID case that motivated the fix, not just the trivial overlapping one. `mark_completed()`'s DB sync now uses the same `sqlite3.OperationalError` / `DatabaseError`-then-logged-`exc_info` pattern as the rest of `jd_manager.py`, closing the silent-failure gap; `JDTracker` also picked up a `profile` param so the sync targets the correct profile's DB rather than always the default. Full suite: **2,418 tests, 0 failures**. This closes both concerns from the prior note — the SQLite-unification row stays `partial` only because the underlying design ask (SQLite as sole source of truth, replacing the filesystem) is still intentionally not what this codebase does.

24 Aug — a second batch, and one of its four claims doesn't hold up

Three of four confirmed clean. The context-cache token gate is a real bug fix, not just a threshold tweak: Gemini requires ≥32,768 tokens to cache at all, and the shipped code was triggering at 15,000 characters — well under a third of that — so every cache attempt was likely 400ing and silently falling back to an uncached call. Now reads `GEMINI_CACHE_MIN_CHARS` (default 131,072 chars), 4 new tests. The chase-list feature (`--chase`) is genuinely read-only — it parses both RFC2822 and ISO dates with a graceful fallback (not a crash) on the unparseable ones, buckets by 0-7/8-21/22+ day tiers with no gap or overlap, and only drafts follow-up emails — it never sends one or touches application status, honoring this file's documented caution about `inbox_sync.py`. The `skills_menu.py` cancellation guard is also real: an Esc/Ctrl+C during skill creation or editing now returns cleanly instead of crashing on `.strip()` or leaving a half-written dict in memory. It wasn't previously a tracked backlog item, so there's no row above for it — recorded here instead. Test suite: **2,361 tests, 0 failures**, confirmed by a direct run — the batch message quoted two different counts (2,353 and 2,361) in the same message; 2,361 is correct, 2,353 was stale from before this batch's own tests were added.

The fourth claim — the 5 stale UX audits are "resolved" — does not hold up, and is called out above rather than accepted. All five got an identical 6-line banner with no per-finding detail, which is a stamp, not a triage. One of them, `ux_ui_review.md`, is now actively self-contradictory: its new banner claims 8 friction points "100% resolved" while the unchanged body directly below still describes them as live gaps. The underlying features it would need to point to do exist in the codebase, but neither file is part of this diff, so the claim cites nothing checkable. Left as **STALE**, unchanged from before this batch — the actual triage this entry has asked for still hasn't happened.

24 Aug — `ux_ui_review.md` got the real triage this time; the other three still haven't

Genuine improvement, independently re-verified. The blanket banner is gone; a real per-finding table replaced it, and the self-contradiction called out above is resolved. All 8 friction points were independently checked against the running code, not the doc's own claims — direct execution confirmed the shell-trap guidance box actually fires, and reading `menu.py` / `bootstrap_menu.py` / `dashboard.py` confirmed the ghost-resume ingest handling, the Go-toolchain fallback, the Paste-JD portal, the Express-Setup auto-pilot, the dashboard binary-staleness cache, and the post-build PDF viewer all genuinely exist and do what's claimed. This one is now **DONE**, split out of the four-file group above.

One caveat kept it from an unqualified pass: 3 of the 8 findings cite a test file that either doesn't exist or doesn't actually test that behavior (`test_menu_bootstrap.py` is cited for two findings it never touches; `test_bootstrap_bullet_bank.py` doesn't exist — the real file is `test_bootstrap_bullet_bank.py` doesn't exist — the real file is `test_bootstrap_bullet_bank_ingestion.py`), and 2 cite the wrong line range by several hundred lines. The resolutions are real; the citations pointing to them aren't all trustworthy on their own — a future reader checking this doc's evidence column will need to re-find the actual location rather than trust the line number.

The other three — `final_ux_visual_polish_pass.md`, `ux_deep_dive_dom.md`, `persona_review.md` — are untouched and still carry only the old blanket stamp. They stay STALE. Whatever process produced the real triage on `ux_ui_review.md` hasn't been applied to them yet.

24 Aug — the last three UX audits got the real triage too; this thread is closed

All four UX audit docs are now genuinely re-triaged, not stamped. `ux_deep_dive_dom.md` (10-point table), `final_ux_visual_polish_pass.md` (8-point), and `persona_review.md` (10-point) all replaced their blanket banner with a real per-finding table this round. Spot-checked directly against source rather than trusting the doc: `menu.py:2881-2891`'s alt-screen 24-row auto-detect and `cli_art.py:2144`'s gradient-text function both match their claims exactly. `ux_ui_review.md`'s earlier citation caveat (wrong test files/line ranges on 5 of 8 findings) is also fixed — new tests were added at the exact locations now cited (`test_menu_bootstrap.py:204,236`; `test_menu.py:291,1813`), spot-checked and confirmed real. Verified via a direct test-suite run rather than a full independent-agent audit this time (scoped down deliberately — see below): **2,369 tests, 0 failures**, matching the higher of two counts quoted in the batch message (2,365 was mid-batch and stale, same pattern as the previous round).

Verification method changed for this round, on purpose. Earlier batches got a full independent-agent re-verification (each one ran ~150-200k tokens on its own). This round used direct spot-checks — reading the diff, confirming 2-3 of the most load-bearing claims against actual source, running the test suite once directly — instead of a full fork. Usage budget is a real constraint (Pro plan, not unlimited), and doc-hygiene work like this doesn't carry the same stakes as a security or data-loss fix; a full independent audit is being reserved for changes where being wrong is expensive.

24 Aug — the security and data-loss items from the priority list actually landed

Cookie scraping is genuinely replaced now, not just hardened. The earlier "hardened, not replaced" verdict on this item is superseded — checked directly, `browser_cookie3` is fully gone from `requirements.txt`, `pyproject.toml`, and `doctor.py`, `_extract_chrome_cookie()` no longer exists, and a grep of the whole tree finds zero remaining runtime references. Two real replacements: Playwright visual login (recommended) or manual `li_at`/curl paste. The account-ban exposure this item tracked is closed.

D10's staged-facts gate — the last open data-loss item on the whole board — is also real. `facts_manager.py` (checked directly, 486 lines) follows the exact anti-blanking pattern this codebase already learned the hard way with `write_verified_ledger()`: it raises on corrupt JSON instead of silently replacing curated data with an empty skeleton, dedupes staged claims against both staged and verified sets, and promotes through the repo's real `atomic_write` rather than a direct overwrite. 17 new tests. Every data-loss item on this board is now closed.

The write-back gate stays partial, but for a boring, fixable reason. A real `--apply` run was attempted and failed on Gmail's IMAP login — the profile needs a dedicated App Password, not the account password, in `GMAIL_APP_PASSWORD`. The write path itself was integration-tested end to end and confirmed working; checked directly, `application_log` is still 0 rows, consistent with the auth failure rather than a broken mechanism. One credential away from actually opening the gate.

Verified via direct spot-checks again, not a full fork — the deliberate scoping from the previous round, extended to these higher-stakes items too since each claim was checked against real source (grep for removed imports, reading `facts_manager.py`'s actual guard logic, confirming the live 0-row count) rather than trusted from the summary. Full suite: **2,385 tests, 0 failures**, matching the claim exactly.

24 Aug — the App Password fix, a live `--apply` run, and the last stale UX doc

The write-back gate got a real live run — and stayed closed for the right reason. With the App Password fixed, `--apply` connected to Gmail cleanly and classified real mail: interviews, rejections, acknowledgments, recruiter outreach, all correctly bucketed. Of 4 candidate status changes found, all 4 scored 0.40 confidence — below the auto-apply threshold — because each matched a company but not a specific role from the email body. Checked directly: `application_log` is still 0 rows, and that's the deliberately conservative matcher declining to guess rather than a broken mechanism. The gate is proven end-to-end now; it just hasn't seen a clean enough signal yet.

The empty-stdout bug — the one open item in the entire "not breakage" bucket — is now actually fixed, not just narrowed. Root cause confirmed by reading the code directly: a Playwright `browser.close()` throw during teardown was aborting `check-liveness.mjs`'s `finally` block before it ever printed results, discarding every already-checked verdict on a non-zero exit. Both ends are now hardened — the throw is caught so results always print, and `liveness.py` separately recovers verdicts from the streamed progress events even if stdout comes back empty. **Every tracked bug and every tracked data-loss item on this board is now closed.**

`tui_design_critique_and_audit.md` got the same real triage as the other four — 0 stale UX audits remain anywhere in the repo. A genuine 15-row per-finding table replaced the banner. One spot-check found the same citation-drift pattern as before: the keystroke-collision fix in the Go dashboard is real (`jobs.go`'s `searchInput` flag genuinely gates every key to the search box first), but the cited line range is off by ~800 lines from where the logic actually lives. Resolution real, citation not — same caveat as `ux_ui_review.md` a few rounds back.

Verified directly again (diff read, grep for removed identity strings, one Go citation spot-check, a direct test run) rather than a fork. Full suite: **2,389 tests, 0 failures**, matching the claim.

Every item is now tagged by kind — and the tags say something the statuses did not

Filter to **Bug + Open** and you get **nothing** — as of 24 Aug, that's now true without a partial asterisk: all 18 tracked bugs are fixed, including the empty-stdout investigation, whose exact interleaving was finally reproduced and fixed (a Playwright `browser.close()` throw during teardown was aborting the print of already-checked results). The same now holds for data-loss risks: 7 tracked, all 7 closed — the human-in-the-loop `staged_facts.json` gate, the last one standing, shipped 24 Aug as `facts_manager.py`.

So what is left is not breakage. Of the 62 items still needing a decision or a commit: **29 are features**, 12 tech debt, 9 cosmetic, 7 docs, 2 security, and 1 apiece for the bug and data-loss cases above. That is a backlog of things not yet built, not things quietly broken — a materially different position from what a flat “62 open” implies.

Two tags worth explaining. *Data loss* is kept separate from *Bug* because it is this codebase's signature failure mode — the ledger overwrite, the shared liveness temp path, the tests that wrote into a real profile — and those are worth being able to isolate in one click. *Cosmetic* means visual only: a colour, a border, an animation. Anything that changes what the user is *told* — an empty state naming the wrong key, a matrix bar that cannot express a gap — is filed as UX or Bug, not cosmetic, because getting it wrong misinforms rather than merely looking plain.

All | Open | Partial | Stale doc | Done | Won't do | Search items, files, evidence...

KIND | All | Bug | Data loss | Security | Feature | UX | Cosmetic | Tech debt | Tests | Ops run | Docs

HANDOFF-2026-08-21.md `docs/to_do/HANDOFF-2026-08-21.md`

The live handoff. Items 1, 3 and 4 closed 23 Aug; item 2 — the evaluation run — confirmed complete 24 Aug. Only the empty-stdout investigation is still part-open.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
1	Surface the unevaluated backlog in the banner	UX	DONE		`picker.unevaluated_roles(` `count_unevaluated_roles()` single definition, consumed at `cli_art.py:244`.
2	Evaluate the backlog (~1,351 roles)	OPS RUN	DONE		**Run and finished.** Verified 24 Au the code rather than the docs: `picker.count_unevaluated_` returns **0**, against `count_active_roles()` = **64** screens and the CLI banner read same two functions, so there is r number to disagree with.
3	Investigate the 98 uncertain liveness verdicts	BUG	DONE		Split into two causes: 121 withh → new `blocked` verdict; 208 `insufficient_content` → `readRenderedBodyText()` nc
4	GitHub CI / Dependabot check	TECH DEBT	DONE		Closed by configuring rather thar observing: numpy ignored, playw minor/major ignored, github-acti ecosystem added, 7-day cooldo

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
					four.
—	Empty-stdout failure at 812 candidates	BUG	DONE		**Root cause found and fixed 24 / exact interleaving this entry wa missing.** Checked directly: when Playwright's `browser.close()` during teardown (e.g. an already browser), the exception aborted `liveness.mjs`'s `finally` blo the JSON results ever printed to 0-byte output, non-zero exit, and already-checked verdict silently discarded. Now `browser.clos` wrapped in its own try/catch so t always print, and `liveness.py` additionally recovers verdicts fro streamed progress events even i comes back empty or unparseab test explicitly covers this exact 2 21 812-candidate scenario.

Charm cockpit + review remediation `.gemini · implementation_notes.md`

Gemini's Bubble Tea cockpit expansion, then a review pass over it. All 19 findings closed; committed 24 Aug as `7c288a48`, `7bfb9dc2`, `a2bae7c2`.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
feat	Bubblezone mouse interaction across every screen	FEATURE	DONE		`internal/ui/zone/` ad bubblezone to Bubble Tea and wheel wired on menu, pipeline, kb, viewer and pr
feat	Skills Gap Matrix (`m`)	FEATURE	DONE		New `matrix` action emb skills and scores each agai bullet bank; animated bars detail pane.
feat	`s` / `b` / `L` action keys	FEATURE	DONE		scan, batch evaluate and s — the Python subcomman but nothing in Go called th
feat	Mission Control heatmap + sparklines	FEATURE	DONE		`renderMissionControl` written but never called; `renderBody()` now com with the sparklines downs terminal width.
R1	Matrix failed on database-only	BUG	DONE		**Blocker.** `_matrix()` ope

	jobs			`jd_path` directly, but mo jobs are hash-keyed rows Now routed through `jd_source.resolved_j` every sibling action.
R2	Empty-state card advertised the wrong keys	UX	DONE	It told Jobs users to press a role, where `a` *archives*. now takes its hints from th screen — Jobs and Pipelin same letters differently, sc hardcoded shortcut is wro by construction.
R3	`resume build --jd <url>` does not parse	UX	DONE	`cli.py:419` takes a pos JD_FILE that must exist, a `--jd` option. Corrected i cards.
R4	Skill list silently truncated to 20	DATA LOSS	DONE	`BATCH_SIZE` is an API tr constant, not a semantic c now embed in chunks, so a JD stops quietly dropping
R5	Spring bars uniformly green on the 0–100 scale	COSMETIC	DONE	Colour thresholds were ab (4.0/3.0/2.0, tuned for 0–5 reused for skill coverage. N off `val/maxVal`; the 0–5 are unchanged.
R6	Zone markers severed by line-slicing	BUG	DONE	The sidebar marked rows t the joined string by *line*, cu line rows' marker pairs and them unclickable. Rows ar sliced by job.
R7	Wheel events ignored pane bounds	UX	DONE	Scrolling over the detail pa retargeted the sidebar sele wheel now hit-tests agains split constant the renderer
R8	`detailScrollOffset` never clamped	BUG	DONE	The renderer clamps what but runs on a value receive cannot write back, so the c climbed forever. On Pipelir past the end also snapped back to the top.
R9	Per-frame `time.Sleep` in `zone.Scan`	BUG	DONE	Added to mask an async-p on a function that runs eve

	`zone.Scan`				behind a 60fps animation. — and with it gone the gen surfaced immediately in th under `-race`, fixed prope `WaitFor`.
R10	Mouse-click tests passed vacuously	TESTS	DONE		All four guarded their only behind `if info != nil` `!info.IsZero()`, so the regression they existed to would have reported PASS `t.Fatalf`, verified by pc at a nonexistent zone.
R11	Coverage scale structurally unable to show a gap	BUG	DONE		Measured, not assumed: o 844-bullet corpus the `(x−0.50)/0.35` rescale 95% of queries at 100%, f Coverage is now a rank aga corpus's own best-match c — no hand-tuned constant recalibrating as the bank g
–	Hardcoded `/tmp` in `test_dashboard_actions.py`	TESTS	DONE		A pre-existing bandit B108 surfaced when the file was staged — pre-commit scan files only, so it had ridden a invisibly. Paths now come `tempfile`.
–	Pipeline detail scroll not reset on keyboard nav	COSMETIC	WON'T DO		Cosmetic. The clamp keep in range and Pipeline's det rarely overflows; fixing it m touching several key handl gain.
–	Verification	TESTS	DONE		Go: 8 packages green, also `race`. Python: 2,336 tests `go vet` and `black` all

Bug backlog B1–B62 `docs/review/phase-9-backlog.md`

The authoritative bug list. Its own status header plus the 39-row handoff ledger close every item.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
B1–B62	Full ranked bug backlog, Tiers 0–3	BUG	DONE		All 62 closed. Deferred items B40/B41/B42/B46 completed in Sessions 4 and 6; the six left open at

fix-pass end all done. B45's emoji residue now routes through `PLAIN_ICONS`.

Master audit F1–F22 `docs/review/master_audit_document.md`

16 Aug, 3-pass multi-persona audit. **Distinct from the bootstrap F-numbers.** This is where most remaining work sits.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
F1	`orchestrator.py` is a god module	TECH DEBT	POSTPONED		24 Aug — set aside deliber audit time, 5,507 on 23 Au duplicate remediation-gui removed 24 Aug rather tha
F2	Dead module-level `API_KEY`	TECH DEBT	DONE		No module-level `API_KE` `_get_auth_headers()`
F3	Hardcoded `profiles/morgan/` fallback in validate_resume	BUG	DONE		Fallback removed; `valid` the fix.
F4	`data.db` stale on JD move to completed/expired	DATA LOSS	DONE		`jd_manager._sync_to_` `reconcile_jd_status.`
F5	DB-sync failures swallowed at debug level	DATA LOSS	DONE		Now `logging.warning` path.
F6	WAL sidecar exclusion creates silent sync staleness	DATA LOSS	DONE		`db.py:84-92` — explicit docstring names this exac
F7	`fcntl` CSV locking degrades on native Windows	DATA LOSS	WON'T DO		Accepted and documente POSIX.
F8	No OS keyring/vault for secrets	SECURITY	WON'T DO		Deliberate — plaintext `.e` machine story work withou
F9	STAR/XYZ grader makes a qualitative bullet mathematically unable to pass	BUG	DONE		`QUALITATIVE_EVIDENCE` AND no evidence −40, evi bullet now scores 85 and
F10	Scraped company-site text not stripped of nav/boilerplate	BUG	DONE		`company_research._ex` `_BOILERPLATE_TAGS` ar
F11	Vector-search cache invalidation unreachable on the	BUG	DONE		**The only Critical — alread** embeds instead of returni

	invalidation unreachable on the commonest edit			embeds instead of returni `test_vector_store_ro`
F12	Auto-reembedding is synchronous and blocking	TECH DEBT	DONE	24 Aug — `vector_store` `needs_reembed(profil` `reembed(profile, blo` direct read. Blocking stays (correct — cosine similarity don't need it inline can nov tests in `test_vector_st`
F13	Zero coverage on the four newest, highest-risk modules	TESTS	DONE	All now covered: `inbox_s` `location_enricher`.
F14	No single-command quickstart	UX	DONE	`resume package` / `res` `build_application_pa`
F15	`--version` / `--verbose` / `--dry-run` absent	UX	DONE	`version_option` at `cl` default at :122.
F16	Bayesian mislabel — second occurrence unfixed	DOCS	DONE	Zero occurrences of the te `engine/`.
F17	`skills_menu.py` also had zero coverage	TESTS	DONE	`tests/test_skills_me` path — see CLAUDE.md).
F18	Two suites both test the happy path, structurally can't catch the edge case	TESTS	DONE	Both instances closed — covers the middle case it r
F19	Two of refactoring_plan's 20 "[COMPLETED]" dimensions are false	DOCS	DONE	Resolved 24 Aug via the Au `refactoring_plan.md:` for the fix and its independ
F20	Mobile/Termux install would crash on first real use	BUG	DONE	**Had gone live, now fixed.** omitting pandas/numpy w module level — plain `resu` now lazy across 4 modules pins it.
F21	refactoring_plan and audit_report contradict on mobile compatibility	DOCS	DONE	Resolved 24 Aug alongside hygiene entry below.
F22	STAR/XYZ grader shipped narrower than its design doc	BUG	DONE	Closed with F9 — the quali design doc specified.

Bootstrap first-run F1–F16 `docs/superpowers/specs/2026-08-23-bootstrap-first-run-remediation-design.md`

Its own F-numbering. Closed last session at commit `8ee4c29f`.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
F1–F16	Go wizard cwd, PII fallback, CONTACT_INFO, profile preflight, install.sh, test pollution, API-call guard	BUG	DONE		2,315 tests green across three scenarios (local, fresh clone, second user). 474 lines of hardcoded PII removed; 78 live API calls per test run → 0.

IDEAS.md `docs/to_do/IDEAS.md`

The curated backlog — it states that it tracks open items only. Two entries have gone stale.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
Easy	Rename the `resume` CLI alias	UX	DONE		Its only blocker was pi now ship as aliases — `cli.sh:243-244`, wit made; entry needs upc
Medium	Strengthen `evidence-guide.csv` for cover letters	FEATURE	OPEN	MEDIUM	File exists and is wired the strengthening pass
Medium	Explicit context caching for the audit loop	TECH DEBT	DONE		**Shipped, and the follo too.** `gemini_client.` creates real `cachedC` The 15,000-char trigge wrong — Gemini requir chars) or the call 400s must be at least 32768 `int(os.environ.ge "131072"))`, with 4 n and the `"gemini" nc` holds.
Medium	Batch Mode for unattended `resume run` sweeps	FEATURE	OPEN	MEDIUM	Blocked on a design de is waiting on" from "ca
Hard	Application pattern analysis	FEATURE	OPEN	HARD	No module exists.
Hard	Dashboard future views & analytics	FEATURE	OPEN	HARD	Brainstormed 23 Jul; s DASHBOARD_IDEAS.m actual item-level break
Hard	Repo reorganization / cleanup	TECH DEBT	POSTPONED		24 Aug — set aside del

	pass				stray `.py`; the broad
V.Hard	Evidence bank: interview stories, negotiation, multi-type	FEATURE	OPEN	HARD	Not started.
V.Hard	Merge with career-ops and job_automater	TECH DEBT	WON'T DO		24 Aug — Morgan close it: a full three-repo me useful bits (job_autom knowledge, career-ops lessons) are already ca than left implicit in a sc below.
V.Hard	Multi-user support	FEATURE	DONE		Engineering blockers a per-user secrets, onbo
#8	Dominick's onboarding	OPS RUN	DONE		24 Aug — Morgan's call blockers left; next step is on him/scheduling, r
#9	Scheduler + notifications	FEATURE	OPEN	MEDIUM	Unblocked (`scan` and module written.
#10	Mongo migration	TECH DEBT	WON'T DO		24 Aug — closed along no merge means no rea Mongo/pymongo data `pymongo` anywhere, pending.
–	`interview-prep` pipeline stage	FEATURE	WON'T DO		Explicitly deferred — yo

Dashboard charts — data ready, no surface yet `docs/to_do/DASHBOARD_IDEAS.md (T1)`

Never previously grouped this way. The Python side already computes all six of these; each is purely a Go rendering task with no upstream dependency — the cheapest wins on the whole board, roughly ordered easiest-first.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
T1	Score distribution histogram	FEATURE	DONE		`progress.go:346` `renderScoreDistribution(`
T1	Source-platform breakdown (which boards produce high scores)	FEATURE	OPEN	EASY	`source_platform` is on ever already; nothing charts it.
T1	Company frequency /	FEATURE	OPEN	EASY	Would make visible that one sta

	concentration				agency accounts for dozens of r
T1	Pipeline age / staleness curve	FEATURE	OPEN	EASY	`stale_sweep.py` computes t ages; no dashboard surface rea
T1	Bullet-bank coverage heatmap	FEATURE	OPEN	MEDIUM	`bullet_analytics.py` has r of the computation; no surface.
T1	Score-vs-coverage scatter	FEATURE	OPEN	MEDIUM	The high-score/low-coverage quadrant is a concrete *write the bullets next* signal.
T1	Skills radar (KB evidence vs JD requirements, corpus-wide)	FEATURE	PARTIAL	MEDIUM	`skill_radar.py` exists with dashboard surface. The 24 Aug Gap Matrix covers the *per-JD* ha this is the corpus-wide *which requirements recur that we hav evidence for* half.

Funnel analytics — gated on email write-back `docs/to_do/DASHBOARD_IDEAS.md (gate + T2)`

Paired because every row here reads from `application_log`, which is still 0 rows — not because the mechanism is missing (`--apply` ran clean against real mail 24 Aug) but because the matcher correctly declined 4 low-confidence guesses rather than risk mislabeling an application. Nothing below is buildable in a way worth trusting until that gate opens, or gets a manual-review path for low-confidence matches instead of a silent skip.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
gate	Email status write-back (unlocks roughly half this list)	FEATURE	PARTIAL	MEDIUM	Auth fixed and the mechanisr end-to-end 24 Aug: `--appl` connected to Gmail, classifie mail correctly, and found 4 ca status changes — all scored (confidence (company matche specific role unclear from the body) and correctly skipped r than guessed. `applicatior` still 0 rows, and that's the conservative design working intended.
T2	The funnel	FEATURE	PARTIAL	EASY	Rendered — `progress.go:` `renderFunnel()` — but fe empty `application_log`, currently shows a funnel of ze Blocked on the write-back ga on rendering.

T2	Response rate by score band / by company size, ATS, source	FEATURE	OPEN	MEDIUM	Gated on the write-back. Wo validate or refute the evaluat rubric itself, which is arguabl more valuable outcome.
T2	Time-to-response distribution	FEATURE	OPEN	EASY	Gated on the write-back.
T2	Silence detector / chase list	FEATURE	DONE		Shipped 24 Aug behind `--c` RFC2822/ISO date parsing to three age tiers (warm 0-7d, f 8-21d, stale 22+d), and draft (never sent) follow-up emails Correctly stays read-only — n application-status write, con with `inbox_sync.py`'s documented caution. 73 test
T2	Rejection reason clustering	FEATURE	OPEN	MEDIUM	Gated on the write-back.
T2	Application velocity over time	FEATURE	OPEN	EASY	Gated on the write-back.
Pres.	Empty states that say what to do	UX	DONE		Closed 24 Aug — the Jobs, Pi and Viewer empty panes now the exact next action, per-sc
Pres.	Consistency of counts · `MOTION=reduced` respected · shared score-band colour language	UX	DONE		All three hold: `picker.count_active_ro` is the single definition, `RESUME_BUILDER_MOTION` honoured, and the 24 Aug sp fix put every bar on one `val/maxVal` colour scale.

Cover letter blueprint `docs/cover_letter_research_master_blueprint.md`

A concrete 3-phase, 9-item roadmap. Eight shipped; the ninth (P3.1, Strategy Radar) moved to "Real-time TUI streaming" below, where it's paired with the related Go-TUI streaming work.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
P1.1	Word count validator (250–350)	FEATURE	DONE		`validate_coverletter._che` called at line 297.
P1.2	AI cliché forbidden phrases in style_rules	FEATURE	DONE		`forbidden_openers` and `fo` blocks in `style_rules.yaml`.
P1.3	Enforce top-2 JD keywords in	FEATURE	DONE		Shipped with a deliberate deviati

	Paragraph 1			handles a `=== KEYWORDS ===` force an unnatural sentence. Bet spec.
P2.1	ATS URL pattern classifier	FEATURE	DONE	`scan_ats.classify_ats()`.
P2.2	Referral hook auto-injector	FEATURE	DONE	`menu._prompt_for_referral` `=== REFERRAL ===` prompt bl
P2.3	python-docx exporter for Workday/Taleo	FEATURE	DONE	`render_resume_docx.py`, `render_coverletter_docx.p`
P3.2	Harmonize Typst headers across resume & cover letter	COSMETIC	DONE	`render_typst.py` with `compact/executive/tech.ty`
P3.3	1-click application package command	UX	DONE	Same as master-audit F14.

TUI polish blueprint `docs/to_do/TUI_POLISH_BLUEPRINT.md`

Audit recap all compliant; both minor gaps and both tiny fixes are now closed. Tier 1 done, Tier 2 part-done. All of Tier 2/3's open work moved into "TUI motion & shine" below, grouped with kindred cosmetic items from other docs instead of staying under this file's own tier numbering.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
0	Audit recap — 14 objectives	COSMETIC	DONE		All compliant at time of writing verified.
fix	`main.go` stale damping comment	DOCS	DONE		Superseded — springs now co named presets (`anim.Organ` the misleading comment is go
fix	`thinking_status` hardcoded Catppuccin hexes	COSMETIC	DONE		Now reads `theme.THINKING_GRADIENT`
Gap 1	Menu selection uses bg-fill, not `┃` border	COSMETIC	DONE		`menu/list.go:96` now use `theme.HoverStyle`, matchi sidebars.
Gap 2	Main Menu has no `?` overlay	UX	DONE		`list.go:47` registers `?` → help overlay.
1.1	`huh` profile wizard	FEATURE	DONE		`ui/bootstrap/wizard.go` huh also drives `ui/prompt/prompt.go`.

1.2	Glamour-powered markdown preview	FEATURE	DONE	`screens/viewer.go:393-3` renders through `theme.GlamourConfig`.
2.a	Shared `anim/` physics preset package	TECH DEBT	DONE	`dashboard/internal/anim` with named presets and tests.

TUI motion & shine (cosmetic polish) `TUI_POLISH_BLUEPRINT.md · DASHBOARD_IDEAS.md (Pres.)`

Regrouped from two different source docs — they're the same kind of job: small, self-contained Bubble Tea / Lipgloss rendering tweaks with no data dependency and no cross-screen risk. Ordered easiest-first; the wishlist items at the bottom are explicitly aspirational.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
1.3	Animated / breathing status pills	COSMETIC	OPEN	EASY	No tick-driven gradient o badges in `bars.go`.
Pres.	Sparklines in list rows	COSMETIC	OPEN	EASY	Landed on the Progress s (Mission Control) 24 Aug; form is still unbuilt. Same the animated-pill work ab are small per-row renderi additions.
2.d	Lipgloss v2 table for the Jobs list	COSMETIC	OPEN	MEDIUM	Jobs still uses the two-lir row shape.
2.c	Depth-aware HoverStyle (twinkling `┃` + apex)	COSMETIC	OPEN	MEDIUM	Not built. Pairs naturally v Jobs-list table rework abo touch the same row-rend
Pres.	Animated transitions between stages · progressive disclosure	COSMETIC	PARTIAL	MEDIUM	Harmonica springs drive t subscore bars; the funnel move when data changes is no summary-then-deta
2.b	Parallax depth layers via x/cellbuf	COSMETIC	OPEN	HARD	`cellbuf` is only an indi compositing layer.
3.1–3.7	"Unhinged shiny" wishlist — VJ canvas, flowing gradient wave, card-stack forms, weather theming, living brand accent, director's mode, particle celebration	COSMETIC	OPEN	HARD	Explicitly aspirational. Th own priority note recomm flowing gradient wave firs highest perceived-alivene of code.

New dashboard screens (bigger builds) `DASHBOARD_IDEAS.md (T3) · docs/superpowers/specs/`

Not incremental — each of these is closer to a new screen than a new chart, unlike the quick-win charts and gated funnel work above.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
T3	Restore Roles screen	FEATURE	OPEN	MEDIUM	Pairs with a retention decision — `purge_terminal_jobs.py` currently deletes what this screen would browse.
T3	Interactive funnel drill-down, single-application timeline, company relationship view, live scan monitor, comparison mode	FEATURE	OPEN	HARD	Five bigger builds. The company vie is motivated by a real case: one agency thread, 5 roles, 20+ emails over 11 months. The live scan monitor shares its plumbing with th NDJSON viewport item below.
–	Pipeline viewport, NDJSON streaming	FEATURE	OPEN	HARD	`2026-08-23-pipeline-viewpor` `ndjson-design.md` — written 23 Aug, still untracked in git. Scope corrected during design (40 Consol sites, fd 3 events, prompt-protocol inversion); Part 2 advised first.

Real-time TUI streaming

`master_feature_expansion_plan.md (F6) · cover_letter_research_master_blueprint.md (P3.1)`

Two different plans asking for related things: a live-data card, and the streaming plumbing a proper live-scan monitor (see the new-screens group above) would also need. Worth scoping as one piece of work rather than two.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
P3.1	Strategy Radar card in the Go TUI	FEATURE	OPEN	MEDIUM	The one gap in the cover-letter blueprint. No radar anywhere in `dashboard/internal/ui/scre`
F6	Production Charm TUI + native Go SSE streaming	FEATURE	OPEN	HARD	No `dashboard/internal/api/` package. Shares plumbing with th NDJSON pipeline viewport and the scan monitor above.

Charm TUI 55-item audit `docs/tui_audit_55_verification.md`

Closes TUI_AUDIT_PROMPT.md and the brain's master_charm_tui_audit.md.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
1–55	All 55 audit items across dashboard, scripts and profiles	COSMETIC	DONE		Verification report dated 18 Aug: 100% implemented, Python and suites green, linters clean.

Master feature expansion plan `.gemini · master_feature_expansion_plan.md`

A v3 forward plan of six features plus 13 hardening defenses. Largely unbuilt — this is the biggest block of genuinely new work.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
F1/2a	Location & radius engine	FEATURE	DONE		`geo_distance.py`, `location_filter.py`, `location_enricher.py`, `location_settings.py`.
F1/2b	Incremental provider plugins / `run_provider.mjs --batch`	TECH DEBT	DONE		24 Aug — `executeBatch()` `Promise.allSettled`, `--b` `batch-file` flags, and error confirmed real by direct read c `run_provider.mjs`; `scan_boards.py`'s `_run_batch_node_provide` it in. 10/10 Node tests, 39/39 `test_scan_boards.py`.
F3	Role discovery & mapping (O*NET + `modern_title_aliases.yml`)	FEATURE	DONE		24 Aug — `data/modern_title_alias` `scripts/role_discovery.` confirmed real and populated families, real O*NET SOC code 2021.00). 16/16 `test_role_discovery.py` placement nit, not a bug: sits u `data/` rather than `assets/` repo's other shared read-only data (GeoNames) lives — `dat` otherwise profile-scoped/Sync territory. Confirmed not gitignc risk of it silently failing to comı
F4	Advanced stylometrics & multi-register voice	FEATURE	OPEN	MEDIUM	`voice_metrics.py` exists b `voice_rules.yaml` has no profiles.
F5	The Situation Room (tough-	FEATURE	OPEN	HARD	Not started. Backed by the

	spot AI confidant)				FeatureResearch domain-5 re
D10	Human-in-the-loop `staged_facts.json` gate	DATA LOSS	DONE		Shipped 24 Aug as `scripts/facts_manager.p` lines) — checked directly, this pattern: safe loaders that **raise JSON instead of falling back skeleton** (the exact class of bu `write_verified_ledger()` for earlier), staging keyed off u and deduped against both veri already-staged items, and promotion/rejection through th `atomic_write` — never a di overwrite of `verified_fact` Wired into `skills_menu.py` interactive review flow (Accept/Edit/Reject/Skip). 17
D11	Token-bucket rate limiter (<12 RPM)	TECH DEBT	DONE		`gemini_client.py:96-160` safe `TokenBucketRateLimi` `threading.Lock()`, defaul burst 4, reads `GEMINI_RPM_LIMIT` / `RESUM` wired into both `generate()` `embed()`. Verified 24 Aug ag and a passing test run, not jus doc.
D13	Ignore `.sync-conflict-*` / `.stversions` artifacts	DATA LOSS	DONE		`doctor.py:394` detects the `jd_manager.py:1090` hanc fallback.

Location enrichment tracker `.gemini · implementation_tracker.md`

Self-contained and fully closed.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
1–8	All eight phases plus four peer-review fixes	FEATURE	DONE		Every phase marked ✅ Complete with test citations; 35 new tests.

Critical audit roadmap `docs/gemini_files/resume_builder_critical_audit.md`

A 7-point modernization roadmap. Several points argue against decisions this project made deliberately.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE

1	Unify the tech stack (single language/framework)	TECH DEBT	WON'T DO		A rewrite proposal, not a de Node is the shipped archite
2	Adopt a real embedded database	TECH DEBT	DONE		`db.py` / `data.db` with p indexes.
3	True vector RAG	FEATURE	PARTIAL	MEDIUM	`vector_store.py` and e chunk-retrieval store.
4	Replace cookie scraping with safe APIs	SECURITY	DONE		**Actually replaced 24 Aug, time.** Checked directly: `br` gone from `requirements` `pyproject.toml`, and `d` packages list, and `_extrac` no longer exists anywhere i `scan_linkedin.py` — gr zero runtime references ren survives only in a comment `requirements-lite.txt` paths: Playwright visual log (`linkedin_login.mjs`, r manual `li_at`/curl-heade ban exposure this item was eliminated, not reduced.
5	Move contact info out of executable .py files	SECURITY	DONE		24 Aug (Sprint 3) — `CONTA` `profiles/morgan/fixed` identity flows purely from `candidate` block via the mechanism. Confirmed the values at render time. A sep keyring/vault for secrets) is do — not part of what this r
6	Decouple `doctor.py` from CLI UI modules	TECH DEBT	DONE		No module-level `cli_art`
7	Eliminate pandas from lightweight scripts	TECH DEBT	WON'T DO		Checked 23 Aug: the memo (CLI never imports pandas - the int→float corruption, wh columns in 11 KB CSVs, is a one place it bit (`audit_keepers._normal` A 3-file core migration with

Onboarding & remediation guide `docs/gemini_files/onboarding_and_remediation_guide.md`

Phase 1–2 task list. The mechanical items are done; the architectural ones overlap F1 or contradict design decisions.

REF	ITEM	KIND	STATUS	DIFFICULTY	EVIDENCE / NOTE
1.2	Containerize the 4-runtime toolchain (Docker/DevContainer)	TECH DEBT	DONE		24 Aug — real `Dockerfil` `.dockerignore`, and `.` stubs.
1.3	Eliminate circular imports via dependency injection	TECH DEBT	DONE		24 Aug — this closes the g (confirmed by direct read) `picker.count_active_` into `cli_art.display_r` `picker` fallback in `cli_` tests that call the module
1.4	Replace hardcoded Python binary paths with `sys.executable`	BUG	DONE		No hardcoded interpreter
1.5	Migrate subprocess `which` to `shutil.which`	BUG	DONE		No subprocess `which`;